

WHEN BEING WRONG IS EXPENSIVE

Four questions that turn “use AI responsibly” into a workflow

An artifact that fails any one of the four is not ready to use.

PRINCIPLE 1

Human-Owned

A named person owns the decision and the outcome.
AI drafts; it cannot be accountable.

ASK Who is accountable if this is wrong?

PRINCIPLE 2

Source-Grounded

The human picks the sources and checks the version.
Grounding cuts fabrication, not risk.

ASK What is this based on?

PRINCIPLE 3

Verifiable

A qualified reviewer can check it against evidence.
Sounding reasonable is not checking.

ASK Can we prove this is correct?

PRINCIPLE 4

Traceable

A reviewer can retrace source, check, and approval.
A finished file is not a record.

ASK Can we reconstruct where this came from?



AI IMPLEMENTATION IN REGULATED ENVIRONMENTS